

Monte Carlo Integration: Direct Lighting

Adam Mally

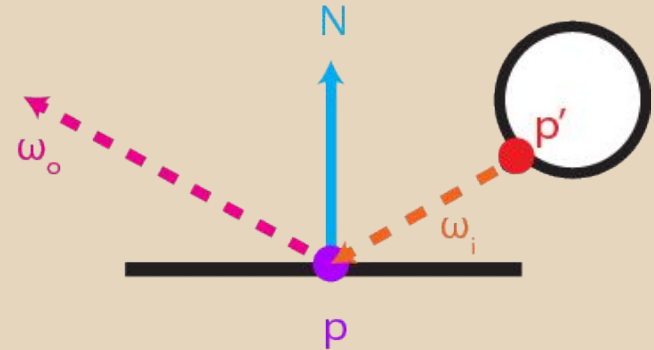
CIS 560 Spring 2015

University of Pennsylvania

Monte Carlo Path Tracer vs Ray Tracer

- MCPT attempts to solve the lighting equation for all visible points in the scene
- Cast many, many rays per pixel in order to gather enough lighting information
- Effects like caustics, soft shadows, anti-aliasing, and depth of field are essentially free
- Without parallelization and other optimizations, converges to a useful image incredibly slowly
- RT casts one ray per pixel (excluding anti-aliasing)
- Makes physically incorrect assumptions about light transportation in order to ensure a fast run time
- Must add special computations to support/fake effects like caustics and depth of field
- Produces a usable image in a relatively short time, even if the image is not physically accurate

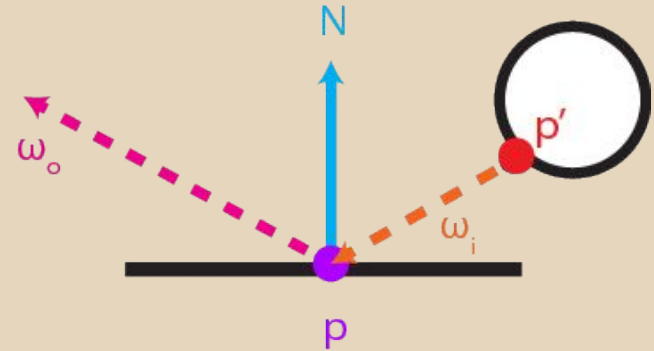
Light Transport Equation - Integral



$$L_o(\mathbf{p}, \omega_o) = L_E(\mathbf{p}, \omega_o) + \int_S f(\mathbf{p}, \omega_o, \omega_i) L_i(\mathbf{p}, \omega_i) V(\mathbf{p}', \mathbf{p}) \cos(\omega_i, \mathbf{N}) d\omega_i$$

- L_o : radiance from a point \mathbf{p} and direction ω_o .
- L_E : emitted radiance from a point (nonzero only if \mathbf{p} is a light source; all other light is *reflected*)
- f : BRDF of surface at \mathbf{p}
- L : radiance at some point \mathbf{x} in direction ω_i .
- V : visibility between \mathbf{p}' and \mathbf{p} (1 when unobstructed, 0 otherwise)
- We incorporate the dot product of ω_i and \mathbf{N} due to Lambert's law
- S is the sphere of all possible ω_i that can reach \mathbf{p}

Light Transport Equation - Sum



$$L_o(\mathbf{p}, \omega_o) = L_E(\mathbf{p}, \omega_o) + \int_{\omega} (f(\mathbf{p}, \omega_o, \omega_i) L_i(\mathbf{p}, \omega_i) V(\mathbf{p}', \mathbf{p}) \cos(\theta_i) d\omega_i)$$

- L_o : radiance from a point \mathbf{p} and direction ω_o .
- L_E : emitted radiance from a point (nonzero only if \mathbf{p} is a light source; all other light is *reflected*)
- f : BRDF of surface at \mathbf{p}
- L : radiance at some point \mathbf{x} in direction ω_i .
- V : visibility between \mathbf{p}' and \mathbf{p} (1 when unobstructed, 0 otherwise)
- We incorporate the dot product of ω_i and \mathbf{N} due to Lambert's law
- For some discrete subset of S , choose n ω to sample and average their reflected light energies
 - The more samples taken, the closer the average is to being physically accurate

Expected Value

- We use Monte Carlo integration to estimate the expected value of the Light Transport Equation (which is an integral)
- The expected value of a function f is the average value of f over some distribution of values $p(x)$ over its domain D
 - $E_p[f(x)] = \int_D f(x)p(x)dx$
 - Note that $p(x)$ must integrate to 1 over the domain D since it is a probability distribution function
 - This means $p(x)$ weights the result of $f(x)$ for any given x while evaluating this integral

Monte Carlo Estimator


- Let's say we want to evaluate some integral $\int_a^b f(x)dx$
- Given a set of **uniform** random variables X_i in the range $[a,b]$, the Monte Carlo estimator says the expected value of our estimator is **equal to the integral of $f(x)$**
- Our estimator is:
 - $F_N = (b - a)/N * \sum_{i=1}^N f(X_i)$
- Whis is this the case?
 - $p(x)$ must equal $1/(b-a)$ since it must be a constant and integrate to 1 over the domain $[a,b]$
 - With algebraic manipulation, we see the following:
 - $E[F_N] = E[(b - a)/N * \sum_{i=1}^N f(X_i)]$
 - $E[F_N] = (b - a)/N * \sum_{i=1}^N E[f(X_i)]$
 - $E[F_N] = (b - a)/N * \sum_{i=1}^N \int_a^b f(x)p(x)dx$
 - $E[F_N] = 1/N * \sum_{i=1}^N \int_a^b f(x)dx$
 - $E[F_N] = \int_a^b f(x)dx$

Monte Carlo Estimator

- What if the set of random variables we use to estimate $E[f(x)]$ is not uniform?
 - e.g. if we're using **importance sampling**
- Our estimator now becomes:
 - $F_N = 1/N \sum_{i=1}^N f(X_i)/p(X_i)$
 - Our only limitation is that $p(X_i)$ be nonzero wherever $f(X_i) > 0$
- Let's prove that this still evaluates to the integral of $f(x)$:
 - $E[F_N] = E[1/N * \sum_{i=1}^N f(X_i)/p(X_i)]$
 - $E[F_N] = 1/N * \sum_{i=1}^N \int_a^b (f(x)/p(x))p(x)dx$
 - $E[F_N] = 1/N * \sum_{i=1}^N \int_a^b f(x)dx$
 - $E[F_N] = \int_a^b f(x)dx$

What is a BRDF?

- A function that evaluates the energy emitted along ray ω_o given a point of intersection p and the direction from which the incoming light emits, ω_i
- Entirely dependent on the attributes of the material sampled at point p
- Examples:
 - A Lambertian BRDF simply takes the material's base color and divides it by π
 - Since we integrate over the entire *hemisphere* of visible directions from p , we need to divide by π so that the integral evaluates to the material's base color
 - A BRDF describing perfect specular reflection would be a discontinuous function that returns 0 (black) for all pairs of ω_o and ω_i that are not reflections of one another around the normal at p



Monte Carlo Path Tracer - Naive Method

- For each pixel, cast N rays through N different points within the pixel
- For each camera ray, get the intersection with the scene
- If the surface is a light source, return its emitted light.
- If the surface is **not** a light source, randomly reflect the camera ray within the hemisphere of possible directions around the point of intersection
 - Recursively trace this ray until you reach some depth limit
 - The recursive tracing is necessary to solve for the $L_i(\mathbf{p}, \omega_i)$ portion of the LTE
- This method is the least biased solution to the Light Transport Equation, but it will also take the longest to converge to a usable image
 - i.e. you'll need to take so many ray samples per pixel that the cost is too prohibitive to use

Monte Carlo Path Tracer - Importance Sampling

- There are several methods we can employ to reduce the number of samples needed to produce a “converged” image
- **Direct Light Sampling:** Light sources are the most important elements in a rendered scene; with no light, there’s no image!
 - For some subset of the rays ω_i , select directions such that each ω_i intersects a given light source at some point
- **BRDF Sampling:** Sample ray directions that have a higher contribution to the color reflected along ω_o
 - Very useful if you have a BRDF with a very narrow set of contributing rays, e.g. a perfectly specular reflective surface

Estimating Direct Lighting

- One method of importance sampling is to estimate how much light travels from an emissive surface and **directly** reaches a given point in the scene and is reflected back along ω_o
- For the purposes of illustration, let's assume there is exactly one light in this scene

$$L_d(\mathbf{p}, \omega_o) = \int_S f(\mathbf{p}, \omega_o, \omega_i) L_d(\mathbf{p}, \omega_i) \text{absdot}(\omega_i, \mathbf{N}) d\omega_i$$

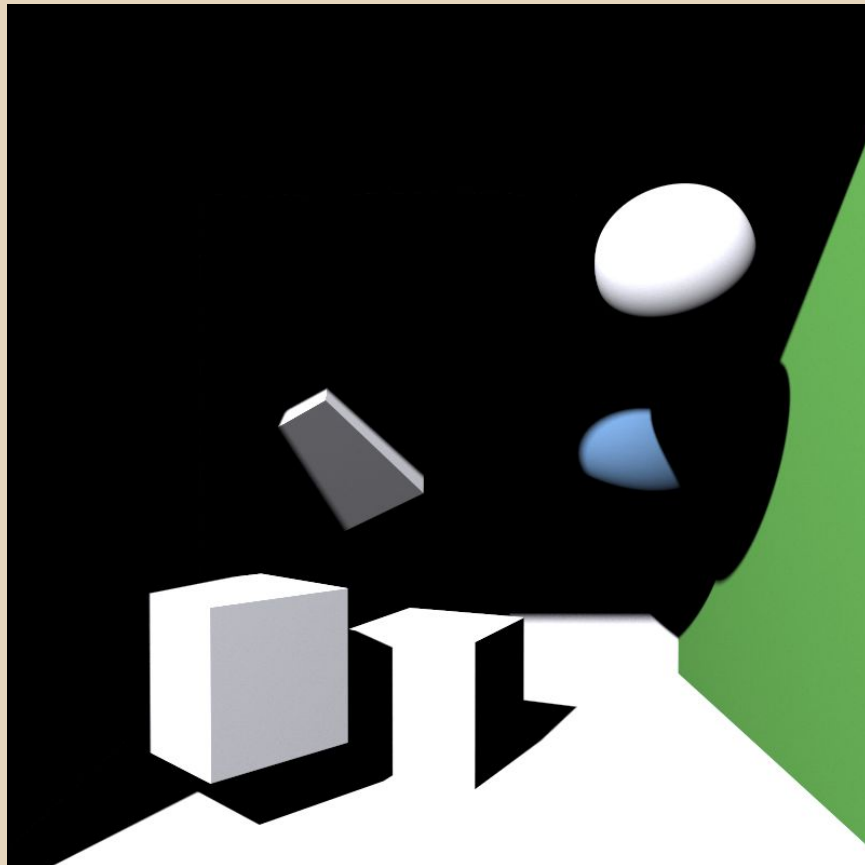
- To write this as an estimate, we need to compute:

$$L_d(\mathbf{p}, \omega_o) = 1/N \sum_j (f(\mathbf{p}, \omega_o, \omega_j) L_d(\mathbf{p}, \omega_j) \text{absdot}(\omega_j, \mathbf{N}) / p(\omega_j))$$

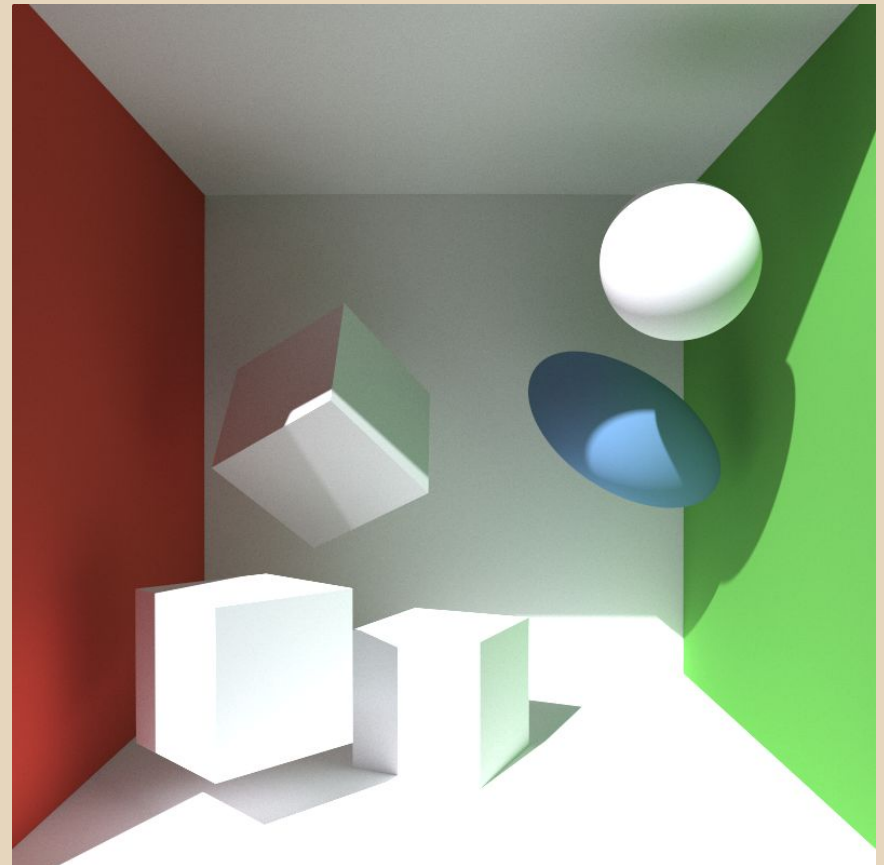
- $p(\omega_j)$ is the probability that ω_j will contribute light
- $f(\mathbf{p}, \omega_o, \omega_j)$ is the BRDF of our point
- $L_d(\mathbf{p}, \omega_j)$ is the light directly reaching \mathbf{p} from direction ω_j
- We want to reduce the variance in our estimates, so we will use importance sampling to choose our various ω_j

For clarity: Direct lighting example

Direct lighting ONLY

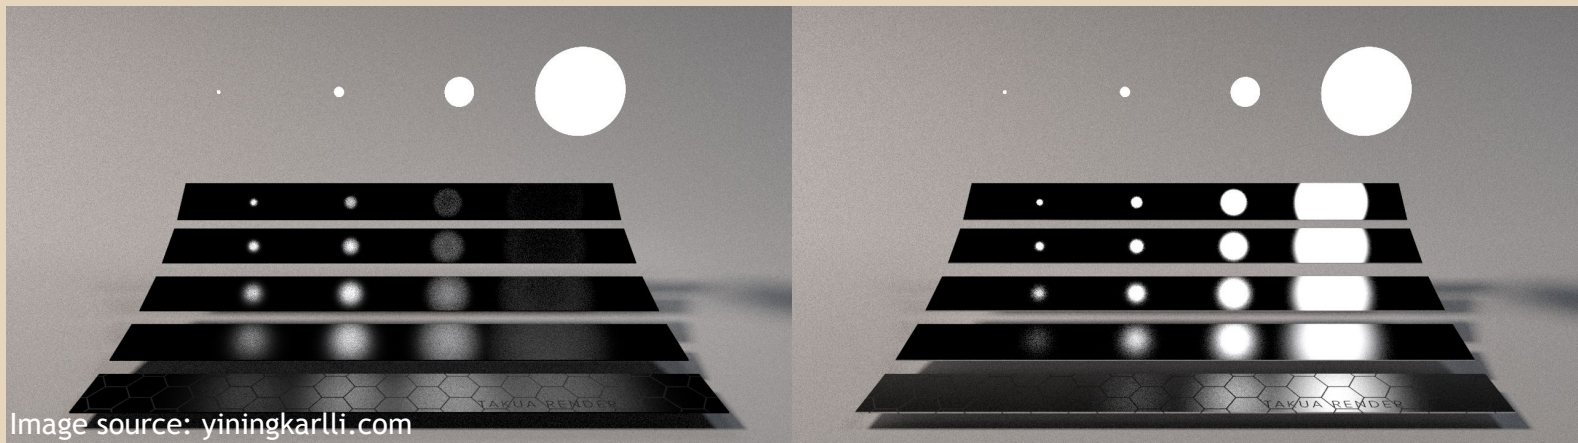


Direct + indirect lighting



Multiple Importance Sampling for DLE

- Since it would be difficult to find a function that matches the PDFs of the light source and the BRDF of our surface point, we will importance sample each PDF separately and weight our results
- Depending on the characteristics of the light and BRDF, one of these sampling methods may be far more effective than the other
- Examples:
 - When a BRDF is more specular, sampling only the **light's PDF** makes it less likely that a large light will contribute to its color (left image)
 - When a BRDF is more diffuse, sampling only the **BRDF's PDF** makes it less likely that small lights will contribute to its color (right image)



Multiple Importance Sampling for DLE

- Since each sampling method is effective in different scenarios, they will produce a high amount of **variance** in our image if not weighted properly
- Since we are sampling two functions (let's call them $f(x)$ and $g(x)$) and have a different importance sampling strategy for each, normally we'd have to choose one strategy over the other
 - Choosing a single strategy will often give poor results for $f(x)$ or $g(x)$
- Simply taking results from our estimator using both sampling methods and averaging the results won't work; variance is *additive*, so it can't be "averaged out"
- So, we need to weight our samples so that they don't contribute too much to our result when the sampling density does not match the shape of our integrand
- Two such weighting methods are known as the **balance heuristic** and the **power heuristic**

MIS: The Balance Heuristic

- If two sampling functions p_f and p_g are used to estimate the value of $\int f(x)g(x)dx$, our Monte Carlo estimator becomes:
 - $1/n_f \sum_{i=1}^{n_f} f(X_i)g(X_i)w_f(X_i) / p_f(X_i) + 1/n_g \sum_{j=1}^{n_g} f(Y_j)g(Y_j)w_g(Y_j) / p_g(Y_j)$
 - n_f is the number of samples taken from the distribution method p_f , and n_g is the number of samples taken from the distribution method p_g
 - w_f and w_g are special weighting functions chosen such that the value of this estimator is the value of the integral $\int f(x)g(x)dx$
- If we use the balance heuristic as w_f and w_g , we get:
 - $w_f = n_g p_g(x) / (n_f p_f(x) + n_g p_g(x))$
 - $w_g = n_f p_f(x) / (n_f p_f(x) + n_g p_g(x))$
- Why does this work? Let's see on the next slide

MIS: The Balance Heuristic

- $w_f = n_{f p_f}(x) / (n_{f p_f}(x) + n_{g p_g}(x))$
- $w_g = n_{g p_g}(x) / (n_{f p_f}(x) + n_{g p_g}(x))$
- Let's consider a case in which the values of $p_f(x)$ and $f(x)$ are low, but $g(x)$ is high
 - With standard importance sampling, $f(x)g(x)/p_f(x)$ will have a high value because $p_f(x)$ is low and $g(x)$ is high
 - With the balance heuristic using $f(x)$, our importance sampling formula becomes:
 - $$\frac{f(x)g(x)n_{f p_f}(x)}{(p_f(x) * (n_{f p_f}(x) + n_{g p_g}(x)))}$$
$$= \frac{f(x)g(x)n_f}{((n_{f p_f}(x) + n_{g p_g}(x)))}$$
- As long as $p_g(x)$ has a distribution that roughly matches $g(x)$, then the denominator won't ever be too small when $g(x)$ is large and $p_f(x)$ is small
- The **power heuristic** is a modification of the balance heuristic that reduces variance even further
- Just square the numerator and both halves of the denominator of the balance heuristic and you've got the power heuristic
- For a proof of why this is more effective, read Eric Veach's thesis

Estimating Direct Lighting: Light PDF Sampling

- Now, back to actually estimating the direct lighting at a point

$$L_d(\mathbf{p}, \omega_o) = 1/N \sum_j (w(\omega_j) f(\mathbf{p}, \omega_o, \omega_j) L_d(\mathbf{p}, \omega_j) \text{absdot}(\omega_j, \mathbf{N}) / p(\omega_j))$$

- Take one sample using the light source's sampling distribution
 - Ask the light source for a random point on its surface \mathbf{p}'
 - Create a ray from \mathbf{p} to \mathbf{p}' . This will be our ω_j
 - Ask the light for its PDF (we'll cover this shortly)
 - Check that the path from \mathbf{p} to \mathbf{p}' is unobstructed (if it is, our received light will be 0)
 - Calculate the light reaching \mathbf{p} along ω_j (this is $L_d(\mathbf{p}, \omega_j)$)
- If the PDF is nonzero and the light emitted along ω_j is nonzero, then we know we can evaluate this function
- Now evaluate $f(\mathbf{p}, \omega_o, \omega_j)$
 - If this BRDF is nonzero, then we can ask \mathbf{p} 's material for its PDF as well so we can use it in our power heuristic
 - We can compute the power heuristic weight now that we have the `light_pdf` and `brdf_pdf`
- Our (first) estimated light is equal to the formula at the top of the page for the ω_j computed using a sample on the light source

Estimating Direct Lighting: BRDF PDF Sampling

- Part two of estimating the direct lighting at a point

$$L_d(\mathbf{p}, \omega_o) = 1/N \sum_j^N (w(\omega_j) f(\mathbf{p}, \omega_o, \omega_j) L_d(\mathbf{p}, \omega_j) \text{absdot}(\omega_j, \mathbf{N}) / p(\omega_j))$$

1. Take one sample using the BRDF's sampling distribution
 - a. Ask p 's BRDF to create a ω_j from its PDF given ω_o
 - b. Store ω_j and the PDF of ω_j
 - c. Sample the BRDF using this ω_j and ω_o
2. If the PDF of the BRDF is nonzero and the light emitted along ω_j is nonzero, then we know we can evaluate this entire function
3. Now evaluate $L_d(\mathbf{p}, \omega_j)$
 - a. If the PDF of our light is nonzero given ω_j , then we want to compute the power heuristic weight using the brdf_pdf and the light_pdf
4. If w_j intersects the chosen light source at all, compute the portion of light from the source that travels along ω_j ; this is $L_d(\mathbf{p}, \omega_j)$
5. Our (second) estimated light is equal to the formula at the top of the page for the ω_j computed using a sample on the BRDF
6. Sum the result of the previous slide with this slide's result for the final direct lighting computation

The PDF of a shape relative to a point

- We need to compute the portion of the hemisphere at p that can see the light source (or, more generally, any shape we want to sample directly)
- We begin by computing the PDF relative to the surface area of the shape
 - $\text{PDF}_{\text{area}} = 1/\text{area}$
- We need to convert this to a PDF relative to the shape's subtended solid angle with respect to p
 - $d\omega_j/dA = \cos(\theta)/r^2$
 - θ is the angle between ω_j and the light's surface normal at the point where ω_j intersects it
 - r^2 is the distance between the point on the shape and p
- We need to divide $1/\text{area}$ by this factor to compute our solid angle PDF:
 - $\text{pdf} = r^2/(\cos(\theta) * \text{area})$
- Let's use an intuitive example for why this works:
 - $d\omega_j = dA \cos(\theta)/r^2$
 - If the surface's normal is parallel to ω_j , then θ is 0. If $r = 1$, then $d\omega_j = dA$, and the equation holds
 - As the surface rotates away from ω_j , its solid angle decreases with $\cos(\theta)$

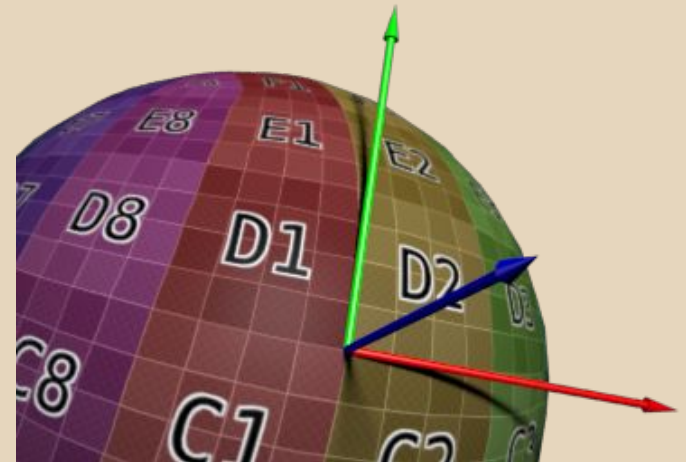
BRDFs: World vs Local Space

- In general, most contemporary path tracers evaluate the output of a BRDF using rays in “normal local space”
 - This treats the surface normal at the input intersection as **always** being $\langle 0 \ 0 \ 1 \rangle$, so the input rays must be transformed accordingly
- Just like when computing the transformation matrix needed to use a normal map, one must calculate the **tangent** and **bitangent** at the point of intersection to make a world \rightarrow normal space matrix
- Let's review the normal map slides...

Normal maps

- To compute the matrix that transforms from texture space to object space, we need three vectors: a **normal**, a **tangent**, and a **bitangent**
 - These form a local orthonormal space
 - The same concept as creating the orientation matrix for a camera
- We already have our **normal**, which is the one we'd use if there were no normal map
- The **tangent** corresponds to our local X axis, so it should align with the U axis of our texture
 - Spherical example: For all points except the very poles of our sphere, we can get the **tangent** by crossing $\langle 0, 1, 0 \rangle$ with our **normal** (remember to normalize the result)
- Likewise, the **bitangent** corresponds to our local Y axis, so it should align with the V axis of our texture
 - Sphere: Get **bitangent** by crossing our **normal** with our **tangent**
- Our transformation matrix can now be created:

$$\begin{vmatrix} \text{T.x} & \text{B.x} & \text{N.x} & 0 \\ \text{T.y} & \text{B.y} & \text{N.y} & 0 \\ \text{T.z} & \text{B.z} & \text{N.z} & 0 \\ 0 & 0 & 0 & 1 \end{vmatrix}$$



Normal maps

- Note that this matrix converts from world to local space, so we need to compute its **inverse**!
- Fortunately, because this is an orthogonal matrix (each column is perpendicular to one another) its inverse is equal to its transpose!
- So, the matrix we want is in fact:

$$\begin{vmatrix} T.x & B.x & N.x & 0 \\ T.y & B.y & N.y & 0 \\ T.z & B.z & N.z & 0 \\ 0 & 0 & 0 & 1 \end{vmatrix}^{-1} = \begin{vmatrix} T.x & T.y & T.z & 0 \\ B.x & B.y & B.z & 0 \\ N.x & N.y & N.z & 0 \\ 0 & 0 & 0 & 1 \end{vmatrix}$$

Tangents and bitangents based on UVs

- In a general case, such as a triangle, we want to be able to compute our tangent and bitangent regardless of UV mapping technique
- If we know three points of a plane on our object and the UVs at those points (which gives us two changes in positions and UVs), we can solve a system of linear equations to compute our **tangent** and **bitangent**:
 - $\Delta Pos1 = \Delta UV1.x * T + \Delta UV1.y * B$
 - $\Delta Pos2 = \Delta UV2.x * T + \Delta UV2.y * B$
- $B = (\Delta Pos2 - \Delta UV2.x * T) / \Delta UV2.y$ OR $(\Delta Pos1 - \Delta UV1.x * T) / \Delta UV1.y$
- $T = (\Delta UV2.y \Delta Pos1 - \Delta UV1.y \Delta Pos2) / (\Delta UV2.y \Delta UV1.x - \Delta UV1.y \Delta UV2.x)$

